Solution 1.1: Setting Up Lab Environment

Access the command line

1. a. If you are doing these exercises from a virtual machine, click on the **Terminal** icon to start a terminal window.

   b. If you have started up an AWS EC2 instance to do exercises, SSH to the instance using its assigned IP address.

2. SSH to the cloud instance using the IP address and key provided.

```
ssh -i <your key filename> ubuntu@<your EC2 ip address>
```

Run cqlsh

3. Start the Cassandra tool 'cqlsh'. You can run 'cqlsh' or any Cassandra / DSE tool from any directory because they are in your '$PATH'.

    ```
    cqlsh
    ```

IMPORTANT: If 'cqlsh' is not able to connect to DataStax Enterprise, use the command 'sudo service dse status' to see if DSE is running. If for some reason DSE is not running, you can start it with 'sudo service dse start'.

In 'cqlsh', note the prompt header changed to `cqlsh>` indicating further commands are for 'cqlsh' instead of the Linux shell.

4. Execute the following query:

    ```
    SELECT * FROM system_schema.keyspaces;
    ```

'cqlsh' displays the results immediately. Notice the built-in keyspaces. Your cqlsh environment is working correctly.

5. Type 'exit' to quit 'cqlsh'.

Solution 1.2: Basic CQL Fundamentals

General setup

Ensure you are in the /labwork/cql directory and you have started 'cqlsh'.

Creating Keyspace and USE

```
CREATE KEYSPACE killrvideo
WITH replication = {'class':'SimpleStrategy','replication_factor': 1};

USE killrvideo;
```

Create videos table

```
CREATE TABLE videos (
  video_id TIMEUUID,
  added_date TIMESTAMP,
  description TEXT,
  title TEXT,
  user_id UUID,
  PRIMARY KEY (video_id)
);
```

COPY data into videos table

```
COPY videos FROM 'videos.csv' WITH HEADER=true;
```

Verify with SELECT

```
SELECT * FROM videos LIMIT 10;
SELECT COUNT(*) FROM videos;
SELECT * FROM videos WHERE video_id = 6c4cffb9-0dc4-1d59-af24-
c960b5fc3652;
```

Remove data from table and exit 'cqlsh'

```
truncate videos;
exit
```

Solution 2.1: Working with Partitions

General setup

Ensure you are in /labwork/partitions/ directory and have started cqlsh. Ensure you have switched to the killrvideo keyspace.

Creating table videos_by_title_year

```
CREATE TABLE videos_by_title_year (
        title TEXT,
        added_year INT,
        added_date TIMESTAMP,
        description TEXT,
        user_id UUID,
        video_id TIMEUUID,
        PRIMARY KEY ((title, added_year))
);
```

Loading data into table

```
COPY videos_by_title_year (title, added_year, added_date,
description, user_id, video_id) FROM 'videos_by_title_year.csv'
WITH HEADER=true;
```

Querying the table

```
SELECT * FROM videos_by_title_year WHERE title = 'Introduction To
Apache Cassandra' AND added_year = 2014;
SELECT * FROM videos_by_title_year WHERE title = 'Sleepy Grumpy
Cat' AND added_year = 2015;
SELECT * FROM videos_by_title_year WHERE title = 'Grumpy Cat:
Slow Motion' AND added_year = 2015;
SELECT * FROM videos_by_title_year WHERE title = 'AzureDev' AND
added_year = 2015;
```

All of these above queries will work for the students.

Querying table with bad partition key

```
SELECT * FROM videos_by_title_year WHERE title = 'Introduction To
Apache Cassandra';
```

```
SELECT * FROM videos_by_title_year WHERE added_year = 2015;
```

InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"

This query fails because of not using the whole partition key.
Since both `title` and `added_year` are part of the partition key, both are required to query.

Solution 2.2: Clustering Columns

General setup

Ensure you are in /labwork/clustering/ directory, that you've started cqlsh and you are using the killrvideo keyspace.

Create upserts

```
CREATE TABLE bad_videos_by_tag_year (
  tag text,
  added_year int,
  added_date timestamp,
  title text,
  description text,
  user_id uuid,
  video_id timeuuid,
  PRIMARY KEY ((video_id))
);

DESCRIBE bad_videos_by_tag_year ;
```

Load bad table

```
COPY bad_videos_by_tag_year (tag, added_year, video_id,
added_date, description, title, user_id) FROM '
videos_by_tag_year.csv' WITH HEADER=true;
```

NOTE: We must explicitly list the column names because this table schema no longer matches the CSV structure.

NOTE: Notice the number of imported rows.

COUNT() the number of rows in the 'bad_videos_by_tag_year'.

```
SELECT COUNT(*)
FROM bad_videos_by_tag_year;
```

Notice the number of rows in the 'bad_videos_by_tag_year' does not match the number of rows imported from 'videos_by_tag_year.csv'. Since 'videos_by_tag_year.csv' duplicates 'video_id' for each unique 'tag' and 'year' per video, Cassandra upserted several records during the COPY. 'video_id' is not a proper partition key for this scenario.

Drop bad table

```
DROP TABLE bad_videos_by_tag_year;
```

Create table videos_by_tag_year

```
CREATE TABLE videos_by_tag_year (
  tag text,
  added_year int,
  video_id timeuuid,
  added_date timestamp,
  description text,
  title text,
  user_id uuid,
  PRIMARY KEY ((tag), added_year, video_id)
) WITH CLUSTERING ORDER BY (added_year desc, video_id asc);
```

Load data

```
COPY videos_by_tag_year FROM 'videos_by_tag_year.csv' WITH
HEADER=true;
```

Query table

```
SELECT * FROM videos_by_tag_year WHERE tag = 'trailer' AND
added_year = 2015;
SELECT * FROM videos_by_tag_year WHERE tag = 'cql' AND added_year
= 2014;
SELECT * FROM videos_by_tag_year WHERE tag = 'spark' AND
added_year = 2014;
```

Count rows

```
SELECT COUNT(*) FROM videos_by_tag_year;
```

NOTE: The number of rows should match the number of rows imported by the COPY command. If not, you had upserts again and will need to adjust your PRIMARY KEY. Ask your instructor for help if necessary.

Query all videos before 2015

```
SELECT * FROM videos_by_tag_year WHERE added_year < 2015;
```

This fails because the partition column `tag` was not queried, which would mean all of the partitions in the table would have to be filtered to get the results for the query.

Solution 2.3: Denormalizing

General setup

Ensure you are in /labwork/denormalization/, you've started cqlsh and you have switched to the killrvideo keyspace.

Creating UDT for encoding

```
CREATE TYPE IF NOT EXISTS video_encoding (
    encoding TEXT,
    height INT,
    width INT,
    bit_rates SET<TEXT>
);
```

Create table videos_by_actor

```
CREATE TABLE videos_by_actor (
  actor TEXT,
  added_date TIMESTAMP,
  video_id TIMEUUID,
  character_name TEXT,
  description TEXT,
  encoding FROZEN<video_encoding>,
  tags SET<TEXT>,
  title TEXT,
  user_id UUID,
  PRIMARY KEY ((actor), added_date, video_id, character_name)
) WITH CLUSTERING ORDER BY (added_date DESC, video_id ASC,
character_name ASC);
```

Load 'videos_by_actor.csv'

```
COPY videos_by_actor
(actor,added_date,video_id,character_name,description,encoding,ta
gs,title,user_id) FROM 'videos_by_actor.csv' WITH HEADER = true;
```

Query videos_by_actor table

```
SELECT * FROM videos_by_actor WHERE actor = 'Leonardo DiCaprio'
LIMIT 10;
```

```
SELECT actor, added_date FROM videos_by_actor WHERE actor =
'Leonardo DiCaprio' LIMIT 10;
```

Create a videos_by_genre table

```
CREATE TABLE videos_by_genre (
  genre TEXT,
  added_date TIMESTAMP,
  video_id TIMEUUID,
  description TEXT,
  encoding FROZEN<video_encoding>,
  tags SET<TEXT>,
  title TEXT,
  user_id UUID,
  PRIMARY KEY ((genre), added_date, video_id)
) WITH CLUSTERING ORDER BY (added_date DESC, video_id ASC);
```

Load 'videos_by_genre.csv'

```
COPY videos_by_genre
(genre,added_date,video_id,description,encoding,tags,title,
user_id) FROM 'videos_by_genre.csv' WITH HEADER = true;
```

Query a genre

```
SELECT * FROM videos_by_genre WHERE genre = 'Musical' LIMIT 10;
```

Solution 3.1: User Defined Types (UDTs)

General setup

Ensure you are in /labwork/udts/, you have started cqlsh and you are in the killrvideo keyspace.

Truncate videos table

```
truncate videos;
```

Alter the 'videos' table to add a 'tags' column

```
ALTER TABLE videos ADD tags SET<TEXT>;
```

Load data

```
COPY videos FROM 'videos.csv' WITH HEADER=true;
```

Remember, we do not need to create the user defined type called 'video_encoding' because we did so in the previous exercise. However, take a look at the code below as a refresher. **Do not run it again or you will get an error!**

```
CREATE TYPE video_encoding (
    bit_rates SET<TEXT>,
    encoding TEXT,
    height INT,
    width INT
);
```

Add an 'encoding' column of the 'video_encoding' type

```
ALTER TABLE videos ADD encoding FROZEN<video_encoding>;
```

Load the data from the 'videos_encoding.csv'

```
COPY videos (video_id, encoding) FROM 'videos_encoding.csv' WITH
HEADER=true;
```

Query first 10 rows

```
SELECT * FROM videos LIMIT 10;
```

Notice the altered table contains data for the new 'tags' and 'encoding' column.

Solution 3.2: Using Counters in CQL

General steps

Ensure you are in /labwork/counters/, you've started cqlsh and you are in the killvideo keyspace.

Create table 'videos_count_by_tag' with a column 'video_count'

```
CREATE TABLE videos_count_by_tag (
    tag TEXT,
    added_year INT,
    video_count counter,
    PRIMARY KEY (tag, added_year)
);
```

Load 'videos_count_by_tag.cql' file

```
SOURCE 'videos_count_by_tag.cql'
```

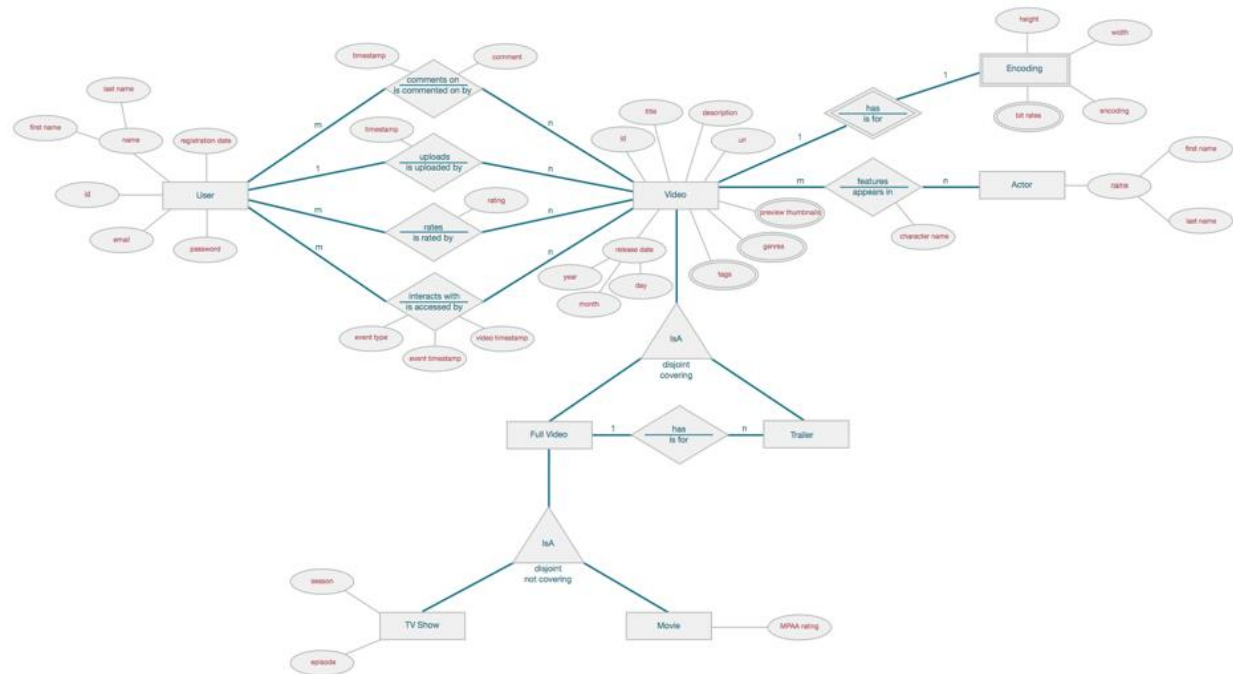**Query table**

```
SELECT * FROM videos_count_by_tag;
```

Update table to increment video

```
UPDATE videos_count_by_tag SET video_count = video_count + 1
WHERE tag = 'cassandra' AND added_year = 2015;

SELECT * FROM videos_count_by_tag WHERE tag = 'cassandra' AND
added_year = 2015;
```
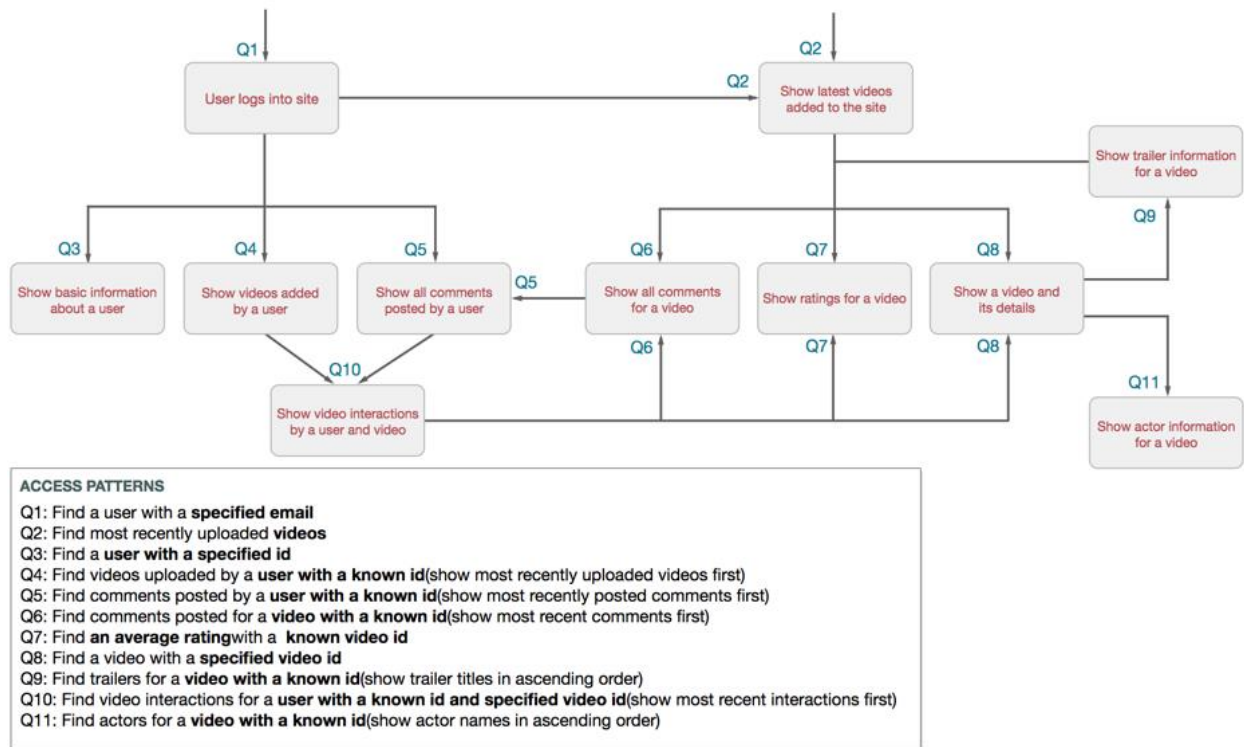
Solution 4.1: Finish a conceptual data model

Diagram KillrVideo entities and their attributes

Solution 4.2: Application Workflow and Access Patterns

Finalized workflow and access patterns

ACCESS PATTERNS

Q1: Find a user with a **specified email**
Q2: Find most recently uploaded **videos**
Q3: Find a **user with a specified id**
Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)
Q5: Find comments posted by a **user with a known id**(show most recently posted comments first)
Q6: Find comments posted for a **video with a known id**(show most recent comments first)
Q7: Find **an average rating**with a **known video id**
Q8: Find a video with a **specified video id**
Q9: Find trailers for a **video with a known id**(show trailer titles in ascending order)
Q10: Find video interactions for a **user with a known id and specified video id**(show most recent interactions first)
Q11: Find actors for a **video with a known id**(show actor names in ascending order)

Solution 4.3: Extend The KillrVideo Logical Model

- **Q1.** Find all user videos that match a specific tag (show the most recent uploaded videos first)

- **Q2.** Find all movies that features a specific actor and release year range (show the most recent videos first, and then sorted by title)

- **Q3.** Find all movies that features a specific actor, genre and release year range (show the most recent videos first, and then sorted by title)

Q1

- What entity type is being stored in a partition or row?

  - Each partition holds the `User posts Video` relationship.

  - Each row represents the video and user information for a specific tag.

- What are the key attribute(s) for this table?

  - The key attribute is `video_id`.

- What attribute is used for the partition key(s) that enables the query?

  - The equality attribute is `tag`.

- What are the clustering column(s) and ordering that support the required results?

  - The clustering column that support ordering is *uploaded_timestamp (DESC)*.

Q2

```
videos_by_actor

actor_name          K
release_year        C↓
title               C↑
video_id            C↑
character_name      C↑
description
type
{genre}
{tags}
```

- What entity or relationship type is being stored in a partition or row?

    o Each partition represents a `Video features Actor` relationship.

    o Each row represents information for a video the actor was featured in.

- What are the key attribute(s) for this table?

    o The key attributes are `actor_name` and `video_id`.

- What attribute is used for the partition key(s) that enabled the query?

    o The equality attribute is `actor_name`.

- What attribute is used for the cluster column(s) that enables the inequality / range scan?

    o The inequality attribute is `release_year`.

- What are the clustering column(s) and ordering that support the required results?

    o The clustering columns that support ordering is *release_year* (DESC) and *title* (ASC).

Q3

```
videos_by_genre_actor

genre               K
actor_name          K
release_year        C↓
title               C↑
video_id            C↑
character_name      C↑
description
type
{genre}
{tags}
```

- What entity or relationship type is being stored in a partition or row?

  - Each partition represents a `Video features Actor` relationship.

  - Each row represents video information for a specific genre and actor.

- What are the key attribute(s) for this table?

  - The key attributes are `actor_name` and `video_id`.

- What attribute is used for the partition key(s) that enabled the query?

  - The equality attributes are `genre` and `actor_name`.

- What attribute is used for the cluster column(s) that enables the inequality / range scan?

  - The inequality attribute is `release_year`.

- What are the clustering column(s) and ordering that support the required results?

  - The clustering columns that support ordering is *release_year* (DESC) and *title* (ASC).

Solution 4.4: Finalizing Physical Data Modeling

General setup

Ensure you are in /labwork/final/ directory, cqlsh is running and you are in the killrvideo keyspace.

Create the tables

Instructor note: The students will be using the slides as guidance to finish filling out the physical data model.

```
CREATE TABLE users_by_email (
    email TEXT,
    password TEXT,
    user_id UUID,
    PRIMARY KEY ((email))
);

CREATE TABLE users (
    user_id UUID,
    email TEXT,
    first_name TEXT,
    last_name TEXT,
    registration_date TIMESTAMP,
    PRIMARY KEY((user_id))
);

CREATE TABLE videos_by_user (
    user_id UUID,
    video_id TIMEUUID,
    title TEXT,
    type TEXT,
    tags SET<TEXT>,
    preview_thumbnails MAP<INT,BLOB>,
    PRIMARY KEY ((user_id), video_id))
WITH CLUSTERING ORDER BY (video_id DESC);

CREATE TABLE comments_by_user (
    user_id UUID,
    posted_timestamp TIMESTAMP,
    video_id TIMEUUID,
    comment TEXT,
    title TEXT,
    type TEXT,
    tags SET<TEXT>,
```

```
        preview_thumbnails MAP<INT,BLOB>,
        PRIMARY KEY ((user_id), posted_timestamp, video_id))
    WITH CLUSTERING ORDER BY (posted_timestamp DESC, video_id ASC);
```

Instructor note: You can suggest taking a look at some of the other tables that already have data types if students are having trouble selecting data types to columns.

Load the schema

```
SOURCE 'killrvideo.cql';
```

Look at schema

```
DESCRIBE KEYSPACE killr_video;
```

Switch to new killr_video keyspace

```
USE killr_video;
```

Copy data into tables from csv files

```
COPY videos FROM 'videos.csv' WITH HEADER=true;

COPY latest_videos FROM 'latest_videos.csv' WITH HEADER=true;

COPY trailers_by_video FROM 'trailers_by_video.csv' WITH
HEADER=true;

COPY actors_by_video FROM 'actors_by_video.csv' WITH HEADER=true;
```

Query to test funtionality

```
SELECT * FROM latest_videos LIMIT 50;
```

Yes, the movie `Gone Girl` exists and was uploaded with the `video_id` 8a657435-0ef2-11e5-91b1-8438355b7e3a.

```
SELECT * FROM videos WHERE video_id = 8a657435-0ef2-11e5-91b1-
8438355b7e3a;
```

The movie was first released on September 26, 2014, and is categorized as *Drama*, *Mystery*, and *Thriller*.

```
SELECT * FROM actors_by_video WHERE video_id = 8a657435-0ef2-
11e5-91b1-8438355b7e3a;
```

The character Desi Collings was played by the actor Neil Patrick Harris.

```sql
SELECT * FROM trailers_by_video WHERE video_id = 8a657435-0ef2-
11e5-91b1-8438355b7e3a;
```

There are three trailers for this movie, and the `trailer_id` of the first one is 8a65751c-0ef2-11e5-9cac-8438355b7e3a.

```sql
SELECT * FROM videos WHERE video_id = 8a65751c-0ef2-11e5-9cac-
8438355b7e3a;
```

Solution 5.1: Creating Secondary Indexes

1. First, the student should create the index using the following syntax:

CREATE INDEX char_name ON actors_by_video(character_name);

2. Second, the student should validate the query using the following queries:

SELECT actor_name FROM actors_by_video WHERE video_id=87c645e8-0ef2-11e5-98f3-8438355b7e3a AND character_name = 'Kelly La Fonda';

SELECT actor_name FROM actors_by_video WHERE character_name = 'George McFly';

Solution 5.2: Creating Materialized Views

1. First, the student should create the materialized view using the following syntax:

```
CREATE MATERIALIZED VIEW videos_by_user AS
SELECT video_id, user_id, title, description, type, url,
       release_date, avg_rating, mpaa_rating, encoding, tags,
preview_thumbnails, genres
FROM videos
WHERE user_id IS NOT NULL AND video_id IS NOT NULL
PRIMARY KEY (user_id, video_id);
```

2. Second, the student should validate the query using the following queries:

```
SELECT video_id, title, user_id FROM videos_by_user WHERE
user_id = 723f6f5f-3658-4449-90d0-4391d63e50a8;
```

Output:

```
video_id                              | title                      | user_id
--------------------------------------+----------------------------+--------------------------------------
7bda7dba-0ef2-11e5-9458-8438355b7e3a  |    Trinity Is Still My Name | 723f6f5f-3658-4449-90d0-4391d63e50a8
7ee22f45-0ef2-11e5-9abc-8438355b7e3a  |        She's Gotta Have It | 723f6f5f-3658-4449-90d0-4391d63e50a8
811caa85-0ef2-11e5-883a-8438355b7e3a  |                       Heat | 723f6f5f-3658-4449-90d0-4391d63e50a8
81552382-0ef2-11e5-be68-8438355b7e3a  |                   Calendar | 723f6f5f-3658-4449-90d0-4391d63e50a8
832f1aa8-0ef2-11e5-a3c4-8438355b7e3a  |                The Coroner | 723f6f5f-3658-4449-90d0-4391d63e50a8
83359ab0-0ef2-11e5-9d31-8438355b7e3a  |              Open Your Eyes | 723f6f5f-3658-4449-90d0-4391d63e50a8
857bc421-0ef2-11e5-830f-8438355b7e3a  |                        5x2 | 723f6f5f-3658-4449-90d0-4391d63e50a8
857e67a3-0ef2-11e5-b31e-8438355b7e3a  | Kekexili: Mountain Patrol | 723f6f5f-3658-4449-90d0-4391d63e50a8
87d20661-0ef2-11e5-8049-8438355b7e3a  |        Strength and Honour | 723f6f5f-3658-4449-90d0-4391d63e50a8
87ea315e-0ef2-11e5-b3fb-8438355b7e3a  |                     Stiffs | 723f6f5f-3658-4449-90d0-4391d63e50a8
88016235-0ef2-11e5-b450-8438355b7e3a  | Bare-Naked Survivor Again | 723f6f5f-3658-4449-90d0-4391d63e50a8
8a53b805-0ef2-11e5-851b-8438355b7e3a  |            The Great Gatsby | 723f6f5f-3658-4449-90d0-4391d63e50a8
8a627fb0-0ef2-11e5-bebb-8438355b7e3a  |              The Lifeguard | 723f6f5f-3658-4449-90d0-4391d63e50a8

(13 rows)
```

Solution 5.3: Implementing Aggregation In Your Data Model

In this exercise, you will:

- Design a logical data model that will be able to store and retrieve aggregated values

Background

During a peer review of an intern's work, your team found that the intern created a schema with the assumption that the queries could use aggregate functions.

It's back to the drawing board, but this time with you working with the hapless intern to come up with a proper design.

- Requirement #1: KillrVideo must be able to count the number of views for each video in a particular month and year. The count does not have to be 100% accurate.

- Requirement #2: KillrVideo must be able to keep track of the total number of uploaded videos, the combined duration (in seconds) of all of these videos, and the average duration (in seconds) of an uploaded video. These statistics must be stored per day.

They do not have to be 100% accurate.

A logical table schema for the `Video` entity has already been created and looks like:

```
videos
video_id                    K
uploaded_timestamp
title
description
type
release_date
{tags}
<preview_thumbnails>
{genres}
```

- Review the above requirement #1.

- The intern made a table to log an entry into whenever a user views a video. The original query would have then used the COUNT aggregate to retrieve the number of views for a particular video.

  o As far as Cassandra is concerned, what is inefficient about the intern's design?

- Come up with an optimal table schema that meets this requirement using Cassandra.

- What would be the query that can retrieve the count of the number of views for a video for a particular year and month?

- Is it possible to retrieve the all-time number of views for a video? If possible, what would be the query?

- Is there anything that needs to be done from the application side?

- How would your design increment a video's view count? Provide the queries and/or write statements needed to do this.

- The intern, being ever helpful, suggested that it would be useful to also display the top 10 videos for each month, based on the number of views.

  - Would it be possible to query this information in Cassandra? Why or why not?

  - If not, is there a way to do this outside of Cassandra using your schema?

- Review the above requirement #2.

- Modify the existing Videos table and/or come up with your own table schema that will meet these requirements.

```
video_statistics

year            K
month           C↓
day             C↓
avg_duration
```

```
video_statistics_counter

year            K
month           C↓
day             C↓
num_uploads     ++
sum_duration    ++
```

- What is the query to retrieve the total number of videos uploaded on a specific day?

  ```
  SELECT num_uploads FROM video_statistics_counter WHERE year = ?
  AND month = ? AND day = ?;
  ```

- What is the query to retrieve the combined duration of all videos uploaded on a specific day?

  ```
  SELECT sum_duration FROM video_statistics_counter WHERE year = ?
  AND month = ? AND day = ?;
  ```

- What is the query to get the average duration of all videos uploaded on a specific day?

```sql
SELECT avg_duration FROM video_statistics WHERE year = ? AND
month = ? AND day = ?;
```

Explain how these values are updated whenever a new video is uploaded.

- The columns in `video-statistics-counter` would be updated first, using the current date:

```sql
UPDATE video-statistics-counter SET num_uploads = num_uploads +
1, sum_duration = sum_duration + ?
WHERE year = ? AND month = ? AND day = ?;
```

- Retrieve the new number of uploaded videos and total duration:

```sql
SELECT num_videos, sum_duration FROM video_statistics_counter
WHERE YEAR = ? AND month = ? AND day = ?;
```

- Calculate the average duration on the application side using the formula `sum_duration` / `num_videos`.
  Store the result in `avg_duration` in `video-statistics`:

```sql
UPDATE video-statistics SET avg_duration = ?
WHERE year = ? AND month = ? AND day = ?;
```

Solution 5.4: Using the CAST function

Steps

1. First consider the following questions:

   What table would you use? Why? ***The videos table. It has the video title and the release date.***

   What filter will you have to apply? ***You will have to filter on 'genres'.***

   How will you filter on the column you decided to use? ***You will need to add a secondary index on 'genres'.***

2. Navigate to the killr_video keyspace.

3. Launch 'cqlsh'.

4. Write the required cql code to satisfy the requirements and validate your work.

   CREATE INDEX idx_genre ON killr_video.videos (values(genres));

   SELECT title, CAST(release_date AS text) AS release_date FROM videos LIMIT 10;

Solution 5.5: Table Optimizations

Given the worst case scenario, the `preview_thumbnails` column would be: 6 hours * 60 minutes per hour * 60 seconds per minute / 20 seconds per image * 20 KB per image = 21,600 KB or ~21 MB

Since there is only one video stored per partition in the `videos` table, these partitions should be manageable.

Using the column size of 21 MB that we just calculated, we can multiply by 500 to get the total size of 10,500 MB or 10 GB.

The size of this partition may make it unmanageable, performance-wise, and would benefit from splitting the partition in some way.

One possibility would be to add a part of uploaded timestamp as the partition key, which will only store in one partition the videos uploaded by the user during a certain time range. Reasonable durations may include all videos uploaded within an hour period, a day, or even a week.

videos_by_user		
user_id	UUID	K
week_first_day	TIMESTAMP	K
video_id	TIMEUUID	C↓
title	TEXT	
type	TEXT	
{tags}	SET<TEXT>	
<preview_thumbnails>	MAP<INT,BLOB>	

The column that may make the most sense would be the `posted_timestamp` column. However since it may be possible for users to post numerous comments in a short amount of time, finding an appropriate duration may be a challenge.

At a minimum, it would be very beneficial to split the table so that
the `preview_thumbnails` column is centralized in a different table. If we take a look at
the `videos` table, it already contains the `preview_thumbnails` column. After splitting
the `comments_by_user` table, a query will first be directed to that table, and after retrieving
the video_id(s), the `video`table can then be queried to retrieve
the `preview_thumbnails` data for those videos.
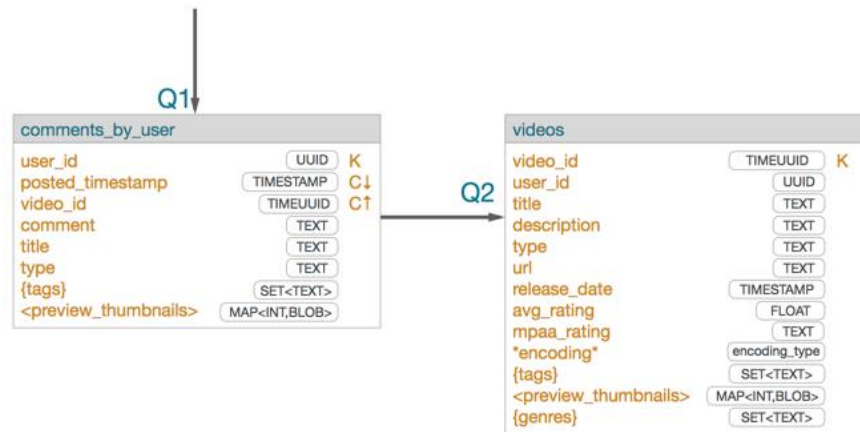
ACCESS PATTERNS

Q1: Find comments posted by a **user with a known id**(show most recently posted comments first)
Q2: Find preview_thumbnails for a **video with a known id**

Solution 6.1: Use Case 1 - Shopping Cart

In this solution we created a shopping_cart table and a product_catalog table. These two tables should satisfy all five access patterns in the application workflow.

```
CREATE TABLE shopping_cart (
     customerid text,
     productid text,
     description text,
     time_added timestamp,
     cost double,
     count int,
     total_cost double,
     PRIMARY KEY ((customerid),productid)
);


CREATE TABLE product_catalog (
     productid text,
     inventory_count int,
     description text,
     cost double,
     PRIMARY KEY (productid)
);
```

Solution 6.2: Use Case 2 - Customer Profile

In this solution, we create types for address, phone and email_address. This allows us to store one row for each customer in the customer table. It also allows us to freely store as many addresses, phone numbers or email addresses as needed for each customer.

```
CREATE TYPE address (
  Address1 text,
  Address2 text,
  city text,
  state text,
  zip_code text
);

CREATE TYPE phone (
  phone_number text
);

CREATE TYPE email_address (
  email_address text
);

CREATE TABLE customer (
    customerid text,
    first_name text,
    last_name text,
    addresses map<text, frozen <address>>,
    phone_numbers map<text, frozen <phone>>,
    email_addresses map<text, frozen <email_address>>,
    user_name text,
    pass_word text,
    PRIMARY KEY (customerid)
);
```

Solution 6.3: Use Case 3 - Sensor Event Tracking

The potential solution below contains a compound partition key including date to avoid large partitions but to allow snapshot time to be queried for a range within that date:

```
CREATE TABLE sensor_data (
    serial_number text,
    date text,
    snapshot_time timestamp,
    facility_id int,
    sensor_type text,
    sensor_value text,
    PRIMARY KEY ((serial_number, date), snapshot_time)
) ;
```