

**DS220**

## 03: Additional Apache Cassandra™ CQL Functionality

# Objectives

Over the next few sections, we will cover:

- Why you should know these other CQL features before you create your data model
- How and why to use collection data types
- How UDTs can be implemented for a more flexible model
- Counter columns and their implications
- How UDFs and UDAs can be used together in your data model

# Introduction

## Other Cool CQL Features

Why do you care? Aren't we talking about data modeling?

- Knowing what other functionality there is can factor into your design
- Apache Cassandra™ has additional data types for flexibility
- These include:
  - Collections
  - Counters
  - UDTs
  - UDFs and UDAs



## Why do I care?

These are tools in your arsenal for modeling data!

- Simplifying table design
- Optimize/streamline table functionality
- Store data more efficiently
- Might change table designs completely

# Collections

# Collection Columns

Group and store data together in a column

- Collection columns are multi-valued columns
- Designed to store a small amount of data
- Retrieved in its entirety
- Cannot nest a collection inside another collection—unless you use FROZEN
  - More on FROZEN to come!



# SET Collection Type

## Creating and inserting with SET

- Typed collection of unique values
- Stored unordered, but retrieved in sorted ordered
- SET example:

```
CREATE TABLE users (
    id text PRIMARY KEY,
    fname text,
    lname text,
    emails set<text>
);

INSERT INTO users (id, fname, lname, emails)
    VALUES ('cass123', 'Cassandra', 'Dev', {'cass@dev.com',
    'cassd@gmail.net'});
```

## LIST Collection Type

Altering a table and updating using LIST

- Like SET—collection of values in same cell
- Do not need to be unique and can be duplicated
- Stored in a particular order
- LIST example:

```
ALTER TABLE users ADD freq_dest list<text>;

UPDATE users SET freq_dest= ['Berlin', 'London',
'Paris']
WHERE id = 'cass123';
```

## Collection Column Syntax, Cont.

### Creating and inserting with MAP

- Typed collection of key-value pairs—name and pair of typed values
- Ordered by unique keys
- MAP example:

```
ALTER TABLE users ADD todo map<timestamp, text>;
```

```
UPDATE users SET todo = { '2018-1-1' : 'create database', '2018-1-2' :  
'load data and test' : '2018-2-1' : 'move to production '}  
WHERE id = 'cass123';
```

## Using Frozen in a Collection

- If you want to nest datatypes, you have to use **FROZEN**
- Using **FROZEN** in a collection will serialize multiple components into a single value
- Values in a **FROZEN** collection are treated like blobs
- Non-frozen types allow updates to individual fields

# UDTs

# User Defined Types (UDTs)

Attach multiple data fields to a column

- User-defined types group related fields of information
- User-defined types (UDTs) can attach multiple data fields, each named and typed, to a single column
- Can be any datatype including collections and other UDTs
- Allows embedding more complex data within a single column

# UDT Example

## Creating the Types

```
CREATE TYPE address (
    street text,
    city text,
    zip_code int,
    phones set<text>
);

CREATE TYPE full_name (
    first_name text,
    last_name text
);
```

## Using Created UDTs

Using the newly created types in a table definition

```
CREATE TABLE users (
    id uuid,
    name frozen <full_name>,
    direct_reports set<frozen <full_name>>,
    addresses map<text, frozen <address>>,
    PRIMARY KEY ((id))
);
```

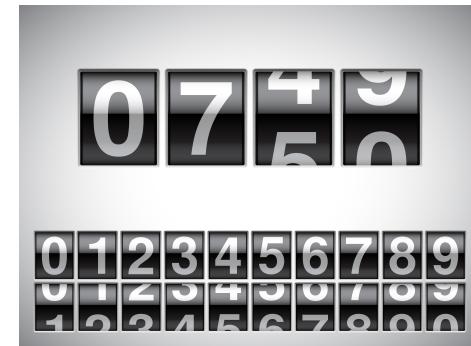
## Exercise 3.1: User Defined Types (UDTs)

# Counters

# Counter Datatype

## What is a counter?

- Column used to store a 64-bit signed integer
- Changed incrementally—incremented or decremented
- Values are changed using UPDATE
- Need specially dedicated tables—can only have primary key and counter columns
  - Can have more than one counter column



## Counter Example

Allow queries for number of videos there are per tag

```
CREATE TABLE moo_counts (
    cow_name text,
    moo_count counter,
    PRIMARY KEY((cow_name))
);

UPDATE moo_counts
SET moo_count = moo_count + 8
WHERE cow_name = 'Betsy';
```

# Counter Considerations

Some things to be aware of

- Distributed system can cause consistency issues with counters in some cases
- Cannot INSERT or assign values—default value is “0”
- Must be only non-primary key column(s)
- Not idempotent
- Must use UPDATE command—DataStax Enterprise rejects USING TIMESTAMP or USING TTL to update counter columns
- Counter columns cannot be indexed or deleted

## Exercise 3.2: Using Counters in CQL

# UDFs and UDAs

# User Defined Functions (UDFs)

## Creating UDFs

- Write custom functions using Java and JavaScript
- Use in SELECT, INSERT, and UPDATE statements
- Functions are only available within the keyspace where it is defined

## Creating UDFs Syntax

- Enabled by changing the following settings in the `cassandra.yaml` file
  - Java: Set `enable_user_defined_functions` to true
  - Javascript and other custom languages:  
Set `enable_scripted_user_defined_functions` to true

## User Defined Aggregates (UDAs)

- DataStax Enterprise allows users to define aggregate functions
- Functions are applied to data stored in a table as part of a query result
- The aggregate function must be created prior to its use in a SELECT statement
- Query must only include the aggregate function itself—no additional columns

## Creating UDF Example

```
CREATE OR REPLACE
    FUNCTION avgState ( state tuple<int,float>, val float )
CALLED ON NULL INPUT
RETURNS tuple<int,float>
LANGUAGE java
AS 'if (val !=null) {
    state.setInt(0, state.getInt(0)+1);
    state.setFloat(1, state.getFloat(1)+val.floatValue());
}
return state;';

CREATE OR REPLACE
    FUNCTION avgFinal ( state tuple<int,float> )
CALLED ON NULL INPUT
RETURNS float
LANGUAGE java
AS 'float r = 0;
if (state.getInt(0) == 0) return null;
r = state.getFloat(1);
r/= state.getInt(0);
return Float.valueOf(r);';
```

## Create UDA Example

```
CREATE AGGREGATE
    IF NOT EXISTS average ( float )
    SFUNC avgState
    STYPE tuple<int,float>
    FINALFUNC avgFinal
    INITCOND (0,0);
```

## Querying with a UDF and a UDA

- The state function is called once for each row
- The value returned by the state function becomes the new state
- After all rows are processed, the optional final function is executed with the last state value as its argument
- Aggregation is performed by the coordinator

```
select average(avg_rating) from videos where release_year = 2002 ALLOW FILTERING;
select average(avg_rating) from videos where title = 'Planet of the Apes' ALLOW FILTERING;
select average(avg_rating) from videos where mpaa_rating = 'G' ALLOW FILTERING;
select average(avg_rating) from videos where genres contains 'Romance' ALLOW FILTERING;
```